

## Deep nets architectures

#### Natalia Khanzhina

Senior Researcher

Machine Learning Lab

nehanzhina@corp.ifmo.ru



#### **Outline**

- MNIST dataset
- **O** CNN
- Autoencoders
- Recurrent Network
- Long-Short Term Memory
- Hopfield network
- Hamming network



#### **MNIST** dataset

- ▼ The training set contains 60000 examples the test set 10000
- Grayscale images of size 28x28
- In past specific features extracted
- examples.http://yann.lecun.com/ex
  db/mnist/



#### Levels of abstraction

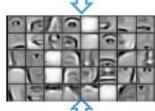
#### **Hierarchical Learning:**

- Natural progression from low level to high level structure as seen in natural complexity
- Easier to monitor what is being learnt and to guide the machine to better subspaces
- A good lower level representation can be used for many distinct tasks

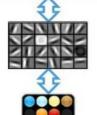
#### Feature representation



3rd layer "Objects"



2nd layer "Object parts"

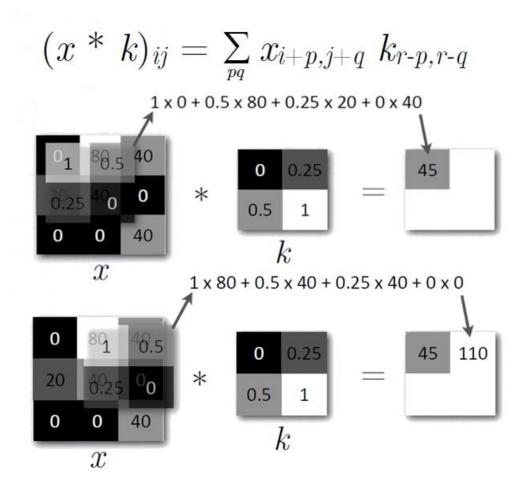


1st layer "Edges"

**Pixels** 



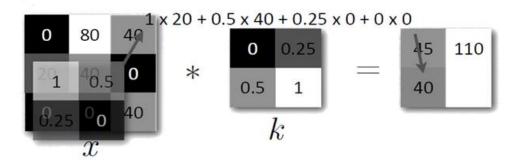
### **Discrete convolution**

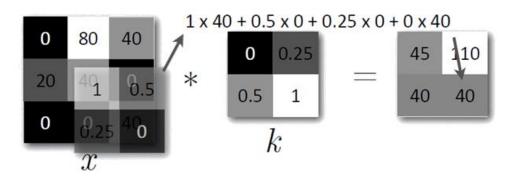




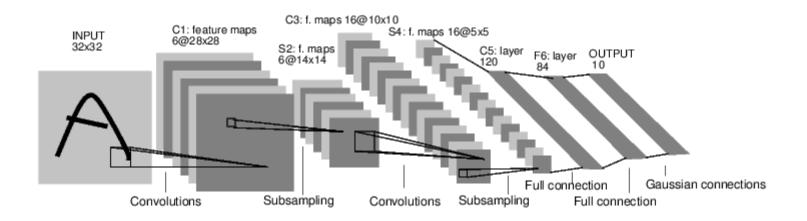
#### Discrete convolution

$$(x * k)_{ij} = \sum\limits_{pq} x_{i+p,j+q} \; k_{r-p,r-q}$$





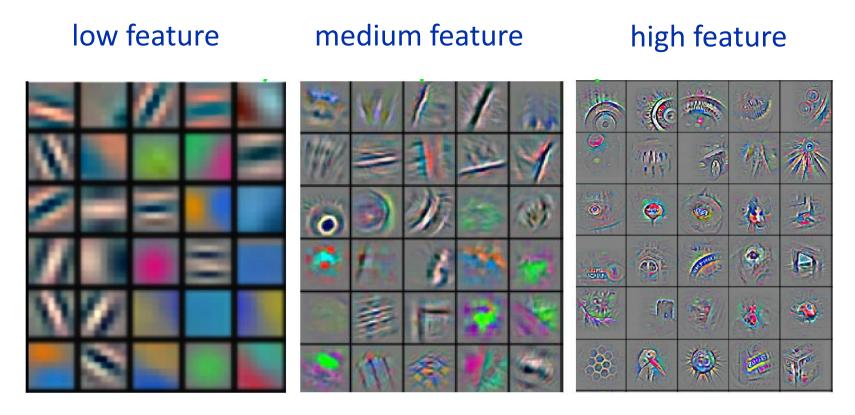
#### **LeNet**



LeCun, Yann, et al. "Gradient-based learning applied to document recognition." *Proceedings of the IEEE* 86.11 (1998): 2278-2324.



#### What do trained kernels look like?



Each kernel composes a local patch of lower-level features into high level representation

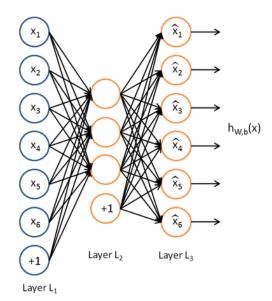


#### **Autoencoders**

Autoencoder: a feed-forward neural network trained to

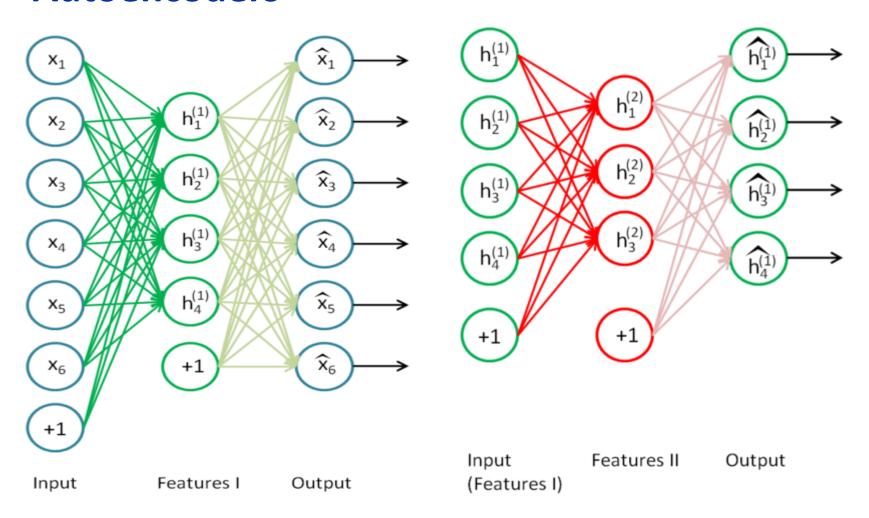
reproduce its input at the output layer

- Make non-linear dimensionality reduction
- Train via backpropagation
- **♥** 1-layer autoencoder is similar with PCA





## **Autoencoders**

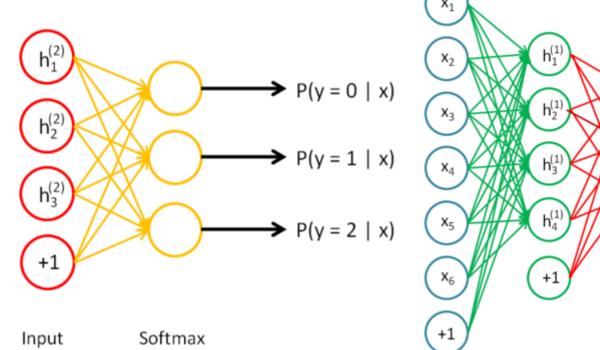


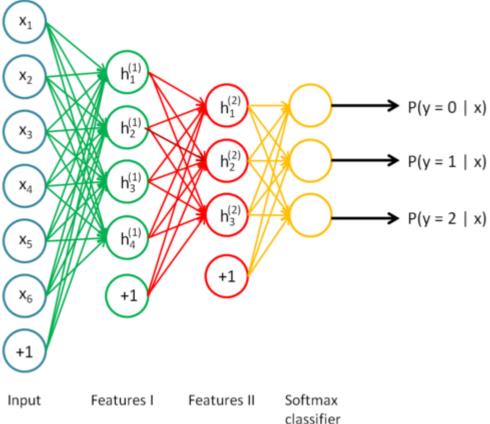


## **Autoencoders**

classifier

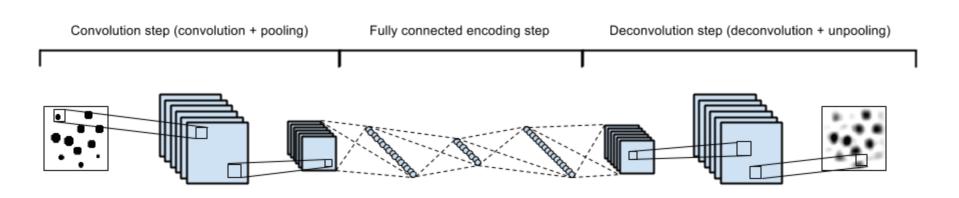
(Features II)





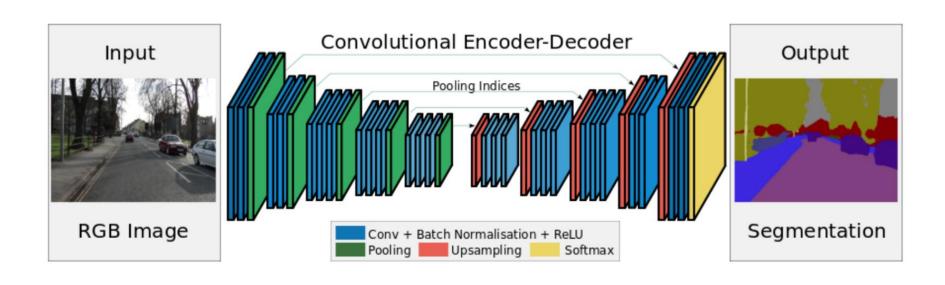


## **C-Autoencoders**





## **Autoencoders for semantic segmentation**



Badrinarayanan, Vijay, Alex Kendall, and Roberto Cipolla. "Segnet: A deep convolutional encoder-decoder architecture for image segmentation." IEEE transactions on pattern analysis and machine intelligence 39.12 (2017): 1481-1495.

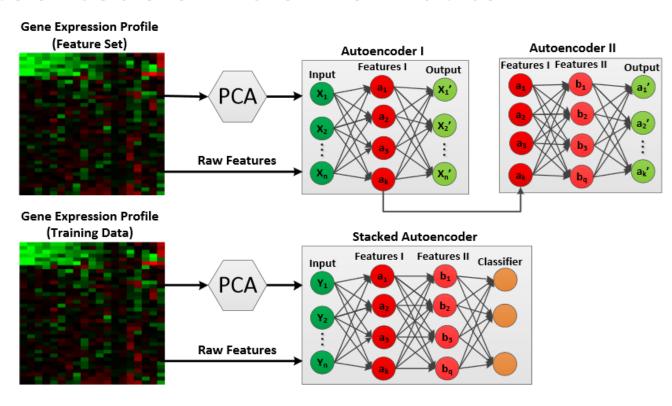


## **Autoencoders for semantic segmentation**

- The task is to classify every pixel
- State-of-the-art results in this field
- The most cool nets are:
- **♥** U-net
- SegNet
- MaskNet



### **Autoencoders in bioinformatics**

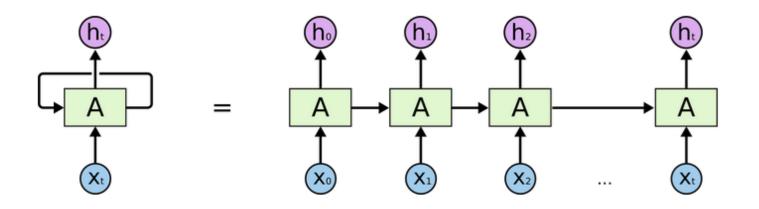


Fakoor, Rasool, et al. "Using deep learning to enhance cancer diagnosis and classification." *Proceedings of the International Conference on Machine Learning*. 2013.

## **Recurrent Neural Network (RNN)**

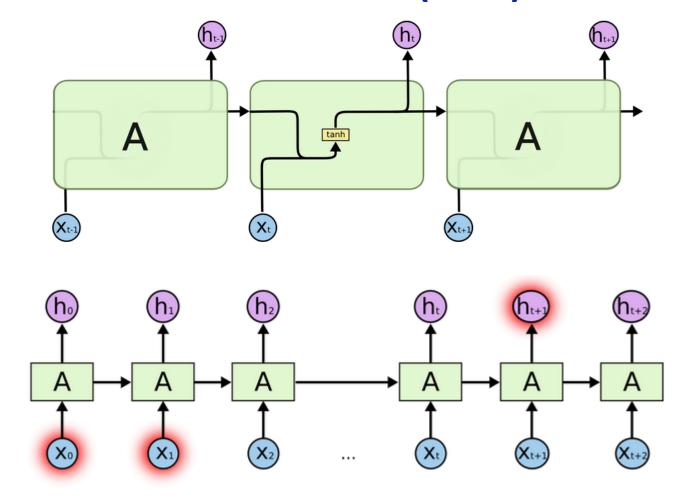
RNN: a neural network with recurrent connections

Suitable for sequence data: time series, text, audio





## **Recurrent Neural Network (RNN)**

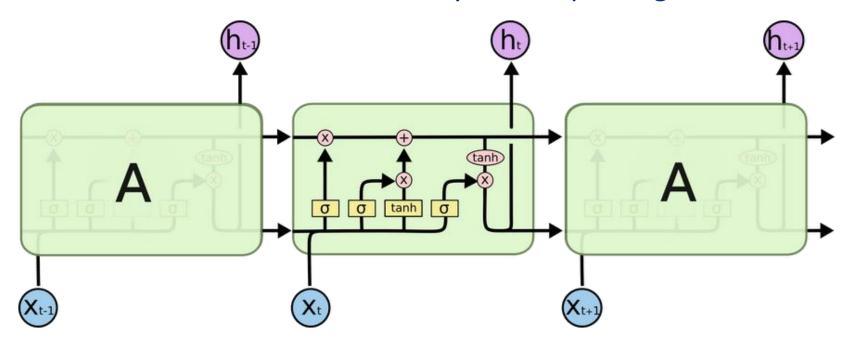




## Long Short-Term Memory (LSTM)

LSTM: a special case of RNN able to learn long-term dependencies

There are four neural network layers in repeating module

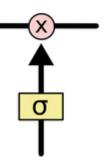




#### LSTM: Cell state

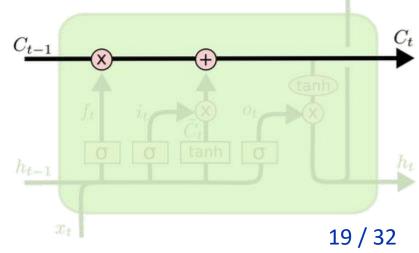
Cell state: runs straight down the entire chain, with only some minor linear interactions

✓ LSTM does have the ability to remove or add information to the cell state, carefully regulated by structures called gates. Gates are a way to sample the input information.



he A

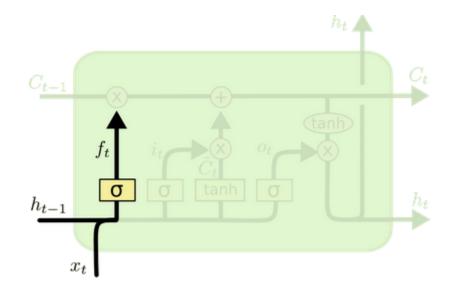
▼ The sigmoid layer outputs numbers between zero and one, describing how much of each component should be let through. LSTM has 3 gates





## LSTM: Forget gate layer

It looks at  $h_{t-1}$  and  $x_t$ , and outputs a number between 0 and 1 for each number in the cell state  $C_{t-1}$ . A 1 represents "completely keep this" while a 0 represents "completely get rid of this"

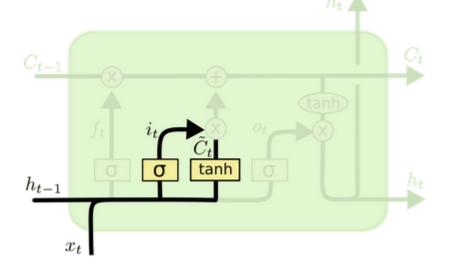


$$f_t = \sigma\left(W_f \cdot [h_{t-1}, x_t] + b_f\right)$$



## LSTM: Input gate layer

The next step is to decide what new information we're going to store in the cell state. This has two parts. First, a sigmoid layer called the "input gate layer" decides which values we'll update. Next, a tanh layer creates a vector of new candidate values,  $C_t$ , that could be added to the state. In the next step, we'll combine these two to create an update to the state

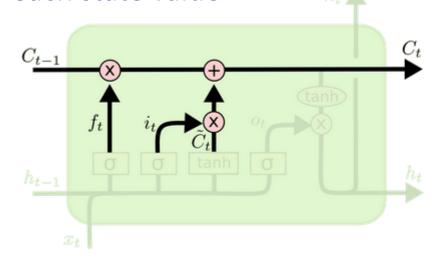


$$i_t = \sigma \left( W_i \cdot [h_{t-1}, x_t] + b_i \right)$$
  
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



## LSTM: Input gate layer

- It's now time to update the old cell state,  $C_{t-1}$ , into the new cell state  $C_t$ . The previous steps already decided what to do, we just need to actually do it
- We multiply the old state by  $f_t$ , forgetting the things we decided to forget earlier. Then we add  $i_t*\mathcal{C}_t$ . This is the new candidate values, scaled by how much we decided to update each state value

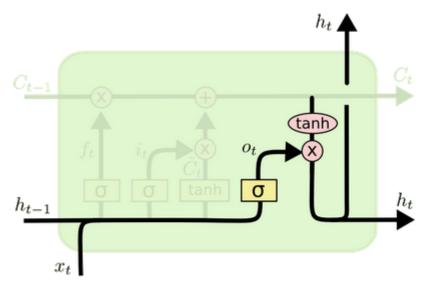


$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$



## LSTM: Output gate layer

▼ The output is based on cell state, but filtered. First, we run a sigmoid layer which decides what parts of the cell state we're going to output. Then, we pass the cell state through tanh (to make the values to be between −1 and 1) and multiply it by the output of the sigmoid gate, so that we only output the parts we decided to.

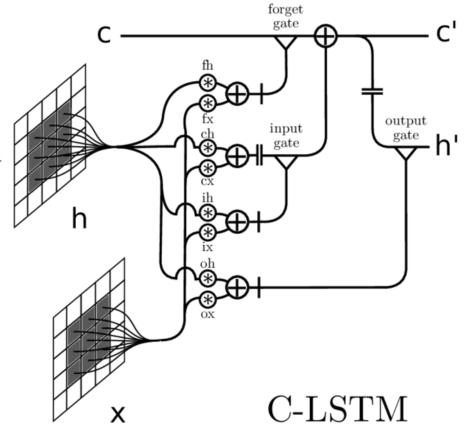


$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh (C_t)$$



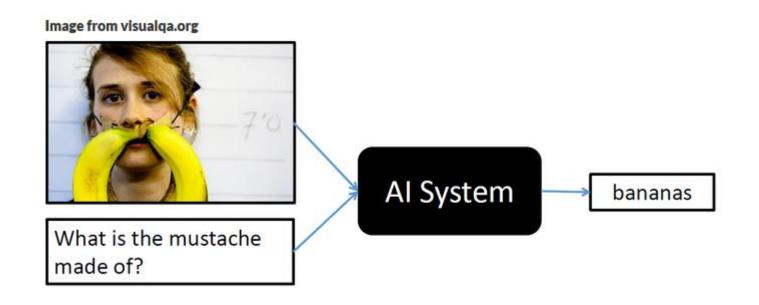
## **C-LSTM**

- ⊕ addition
- + sigmoid
- **→** tanh
- **-**⊛- convolution
- $\mathbf{Y}$  gating



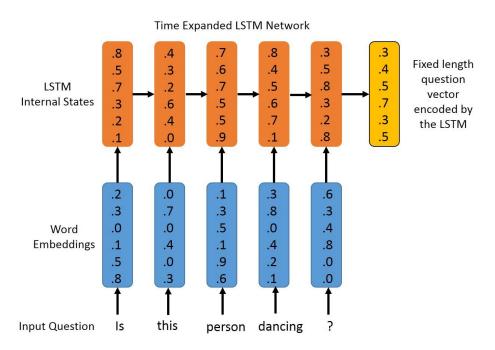


## Visual question answering





## Visual question answering: LSTM



Model	Accuracy
BOW+CNN	44.30%
LSTM-Language only	42.51%
LSTM+CNN	47.80%

26/32

## Hopfield neural network

- Recurrent neural network
- ▼ John Hopfield, 1982
- ▼ Has "associative" memory
- Guaranteed to converge to a local minimum
- 1 layer
- Fast training



## Hopfield neural network

Initial weights: 
$$w_{ij} = \begin{cases} \sum_{k=0}^{m-1} x_i^k x_j^k, i^l j \\ 0, i = j \end{cases}$$

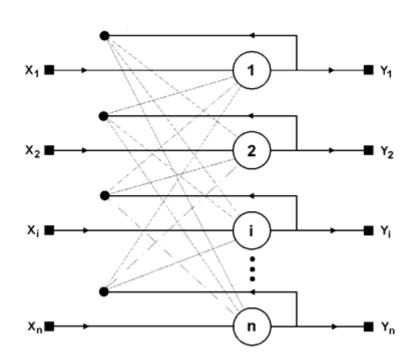
- Predict:
- An unknown signal is sent to the inputs of the network. Setting axon values:

$$y_i(0) = x_i, i = 0...n-1,$$

2. The new state of neurons and axons is calculated:  $s_j(p+1) = \sum_{i=0}^{n-1} w_{ij} y_i(p)$ 

$$y_j(p+1) = f[s_j(p+1)]$$

3. Check if the output values of the axons have changed during the last iteration. If yes - go to point 2, otherwise the end



## Hamming neural network

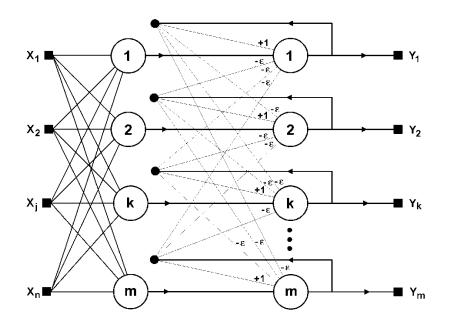
- Hamming distance from the tested image to all samples
- 2 layers
- V Initial weights:  $w_{ik} = \frac{x_i^k}{2}$ , i=0...n-1, k=0...m-1 $T_k = n/2$ , k=0...m-1
- Predict:
- 1. An unknown signal is sent to the inputs of the network. Setting axon values:

$$y_j^{(1)} = s_j^{(l)} = \sum_{i=0}^{n-1} w_{ij} x_i + T_j$$
  $y_j^{(2)} = y_j^{(1)}, j = 0...m-1$ 

2. The new state of neurons and axons is calculated:  $s_j^{(2)}(p+1) = y_j(p) - e^{\sum_{k=0}^{m-1} y_k^{(2)}(p), k^l j}, j = 0...m-1$ 

$$y_j^{(2)}(p+1) = f[s_j^{(2)}(p+1)] j = 0...m-1$$

3. Check if the output values of the axons have changed during the last iteration. If yes - go to point 2, otherwise the end





#### **Materials**

#### Presentation was prepared using:

- 1. <a href="http://avisingh599.github.io/deeplearning/visual-qa/">http://avisingh599.github.io/deeplearning/visual-qa/</a>
- 2. http://colah.github.io/posts/2015-08-Understanding-LSTMs/

#### Task

- 1. Implement Hopfield or Hamming network on Python
- 2. Train it on MNIST dataset (~0.15n samples)



# Thanks for attention! Questions?